Introduction to OpenACC directives

Carlos Junqueira-Junior
Science Computing Researcher Engineer

DynFluid
Fluid Dynamics Laboratory

Arts et Métiers
Sciences et Technologies

Summary

Summary

Summary

Introduction

IDRIS/CNRS - FRANCE

- The slides are based on the material provided by the Institute for *Development and Resources in Intensive Scientific Computing* (IDRIS/CNRS);
- One can find more information on the ▸ IDRIS website .

Introduction

The Ways to GPU

Programming Effort and Technical Expertise

Libraries	Directives	Programming Languages
– cuBLAS – cuSPARSE – cuRAND – AmgX – MAGMA	– OpenACC – OpenMP 5.0	– CUDA – OpenCL
– Min. change in the code – Max. performance	– Portable, simple, low intrusiviness – Efficient	– Complete rewriting, complex – Non portable – Exellent performance

Introduction

Brief History

OpenACC

- https://www.openacc.org
- AMD, CRAY, NVIDIA, Oak Ridge
- First standard 1.0 (11/2011)
- Last standard 3.0 (11/2019)
- Main compilers:
    - PGI;
    - CRAY (only CRAY hardware);
    - GNU ($\geq$ 5.7).

OpenMP

- https://www.openmp.org
- First standard 4.5 (11/2015)
- Last standard 5.0 (11/2018)
- Main compilers:
    - PGI;
    - CRAY (only CRAY hardware);
    - CLANG;
    - IBM XL;
    - GNU ($\geq$ 7.0).

Introduction

Definitions

- Gang(OpenACC)/Teams(OpenMP) : Coarse-grain parallelism;
- Worker(OpenACC) : Fine-grain parallelism;
- Vector: Group of threads executing the same instruction (SIMT);
- Thread : Execution entity;
- SIMT : Single Instruction Multiple Threads;
- Device : Accelerator on which execution can be offloaded (ex: GPU);
- Host : Machine hosting 1 or more accelerators and in charge of execution control;
- Kernel : Piece of code that runs on an accelerator;
- Execution thread : Sequence of kernels to be executed on an accelerator.

Introduction

CPU and GPU Architectures

CPU

- Several Arithmetic Logic Units (ALU);
- Control unit to control the ALUs;
- Fast cache memory;
- Dynamic Random Access Memory;
- High capacity memory;
- $\approx 1/8$ of the area is dedicated to the computing.
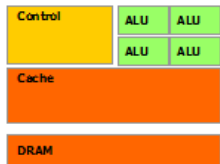
GPU

- Hundreds of (ALU), grouped in several multiprocessors;
- Several control units;
- Several cache memories;
- One Dynamic Random Access Memory;
- High bandwidth memory;
- $\approx 3/4$ of the area is dedicated to the computing.

GPUs are designed such that more transistors are devoted to data processing rather than data caching and flow control as CPUs.

Introduction

Execution Model

Kernels, data transfers and memory allocation are managed by the host (CPU).

Programming GPU-Accelerated Systems
Separate CPU System and GPU Memories

GPU Developer View

Introduction

Execution Model

Kernels, data transfers and memory allocation are managed by the host (CPU).

Introduction

Execution Model

Three levels of parallelism:

- Coarse grain: gang;
- Fine grain: worker;
- Vectorization: vector;
- Sequential: seq.

Introduction

Execution Model

Three levels of parallelism:

- Coarse grain: gang;
- Fine grain: worker;
- Vectorization: vector;
- **Sequential: seq.**

Introduction

Execution Model

Three levels of parallelism:

- **Coarse grain: gang;**
- Fine grain: worker;
- Vectorization: vector;
- Sequential: seq.

Introduction

Execution Model

Three levels of parallelism:

- **Coarse grain: gang;**
- **Fine grain: worker;**
- Vectorization: vector;
- Sequential: seq.

Introduction

Execution Model

Three levels of parallelism:

- **Coarse grain: gang;**
- **Fine grain: worker;**
- **Vectorization: vector;**
- Sequential: seq.

Introduction

Parallelism Models

- Gang-redundant (GR): All gangs run the same instructions redundantly;
- Gang-partitioned (GP): Work is shared between gangs (!$ACC loop gang);
- Worker-single (WS): One worker is active in GR or GP mode;
- Worker-partitioned (WP): Work is shared among workers of a gang;
- Vector-single (VS): One vector channel is active;
- Vector-partitioned (VP): Several vector channels are active.

These models must be combined in order to get the best performance from the calculator.

Introduction

Parallelism Models

- The execution of a kernel uses a set of threads that are mapped on the hardware resources of the accelerator;
- Threads are grouped within team of the same size, with one master thread per team (gang definition);
- Each team is spread on a 2-D thread-grid (worker-vector);
- One worker is actually a vector of $vector_{length}$ threads;
- The total number of threads is:

$$nb_{threads} = nb_{gangs} * nb_{workers} * vector_{length} \, .$$

Important Notes

- There is no synchronization among gangs;
- The compiler can decide to synchronize the threads of a gang (all or part of them);
- Threads of a worker works in SIMD* (single instruction multiple data) fashion.

All threads run the same instruction at the same time, for example on NVidia GPUs, groups of 32 threads are formed.

Introduction

NVIDIA P100 Restrictions

- Number of gangs is restricted to $2^{31} - 1$;
- The 2-D thread-grid size ($nb_{workers} * vector_{length}$) is limited to 1024;
- Due to register limitations the size of the grid should be less than or equal to 256 if the programmer wants to be sure that a kernel can be launched;
- The size of a worker ($vector_{length}$) should be multiple of 32;
- **PGI limitation:** In a kernel that contains calls to external subroutines (not seq), the size of a worker is set at to 32.

Introduction

Porting Strategy

1. Identify the compute intensive loops;
2. Add OpenACC directives;
3. Optimize data transfers and loops;
4. Repeats 1 to 3 until everything is on the devices.

Introduction

PGI Compiler

Information

- Founded in 1989;
- Acquired by NVIDIA in 2013;
- Develops compilers, debuggers and profilers.

Compilers

- Latest version: 19.10;
- Hands-on version: 19.10;
- C: pgcc;
- C++: pgc++;
- Fortran: pgf90, pgfortran.

Activate OpenACC

- -**acc** : Activates openACC support;
- -**ta**=<**options**> : OpenACC options;
- -**Minfo=accel** : Display informations about compilation. The compiler will do implicit operations that are important to the developers.
  Highly recommended to use!

Tools

- nvprof: CPU/GPU profilers;
- nvvp: nvprof GUI.

Introduction

"–ta" Options

Each GPU generation presents different capabilities.
For NVIDIA hardware, it is represented by a number:

- K80: cc35;
- P100: cc60;
- V100: cc70.

A compilation using V100 features is given by:

- **-ta=tesla:cc70**

Complete documentation: ▸ OpenACC User-Guide

Memory Management

- Pinned: The memory location on the host is pinned. It might improve data transfers;
- Managed: The memory of both, host and device, are united.

Introduction

Information available with -Minfo=accel

```fortran
4    program loop
5      integer :: a(10000)
6      integer :: i
7
8      !$ACC parallel loop
9      do i=1,10000
10       a(i) = i
11     enddo
12
13   end program loop
```

pgfortran -O0 -acc -ta=tesla -Minfo=accel loop.f90

```
loop:
      8, Generating Tesla code
         9, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
      8, Generating implicit copyout(a(:)) [if not already present]
```

Introduction

Information available with -Minfo=accel

```
15      !$acc parallel loop
16      do i=1,10000
17        a(i) = i
18      enddo
19      !$acc parallel loop reduction(+:summ)
20      do i=1,10000
21        summ = summ + a(i)
22      enddo
23    end program reduction
```

pgfortran -O0 -acc -ta=tesla -Minfo=accel reduction.f90

```
reduction:
      15, Generating Tesla code
          16, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
      15, Generating implicit copyout(a(:)) [if not already present]
      19, Generating Tesla code
          19, Generating reduction(+:summ)
          20, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
      19, Generating implicit copyin(a(:)) [if not already present]
          Generating implicit copy(summ) [if not already present]
```

Introduction

Information available with -Minfo=accel

```
15        !$acc parallel loop
16      do i=1,10000
17        a(i) = i
18      enddo
19        !$acc parallel loop reduction(+:summ)
20      do i=1,10000
21        summ = summ + a(i)
22      enddo
23    end program reduction
```

pgfortran -O0 -acc -ta=tesla -Minfo=accel reduction.f90

```
reduction:
     15, Generating Tesla code
         16, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
     15, Generating implicit copyout(a(:)) [if not already present]
     19, Generating Tesla code
         19, Generating reduction(+:summ)
         20, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
     19, Generating implicit copyin(a(:)) [if not already present]
         Generating implicit copy(summ) [if not already present]
```

Introduction

Data region in order to avoid CPU-GPU communication

```
15      !$acc data create (a(1:10000))
16      !$acc parallel loop
17      do i=1,10000
18        a(i) = i
19      enddo
20      !$acc parallel loop reduction(+:summ)
21      do i=1,10000
22        summ = summ + a(i)
23      enddo
24      !$acc end data
```

pgfortran -acc -ta=tesla -Minfo=accel reduction_data_region.f90

```
reduction:
     15, Generating create(a(:)) [if not already present]
     16, Generating Tesla code
         17, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
     20, Generating Tesla code
         20, Generating reduction(+:summ)
         21, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
     20, Generating implicit copy(summ) [if not already present]
```

Introduction

Code Profiling

PGI_ACC_TIME

- Command line tool;
- Environment variable for PGI compilers (PGI_ACC_TIME=1);
- Provides basic information such as:
  - Time spent in kernels;
  - Time spent in data transfers;
  - How many times a given kernel is executed;
  - The number of gangs, workers and vector size mapped to hardware.

nvprof

- Command line tools;
- Options that can give you a fine view of the code.

nvvp, pgprof and NSight Graphics Profiler

- Graphical interface for nvprof.

Introduction

PGI_ACC_TIME=1 — parallel-data-multi.f90

```
29      !$ACC parallel copyout( a(:s) )
30      !$ACC loop
31      do i=1,s
32        a(i) = i
33      enddo
34      !$ACC end parallel
35
36      do j=1,p
37        !$ACC parallel copy( a(:s) )
38        !$ACC loop
39        do i=1,s
40          a(i) = a(i) + 1
41        enddo
42        !$ACC end parallel
43      enddo
```

```
para  NVIDIA  devicenum=0
  time(us): 29,740
  29: compute region reached 1 time
      29: kernel launched 1 time
          grid: [79]  block: [128]
           device time(us): total=4 max=4 min=4 avg=4
           elapsed time(us): total=368 max=368 min=368 avg=368
  29: data region reached 2 times
      34: data copyout transfers: 1
           device time(us): total=22 max=22 min=22 avg=22
  37: compute region reached 1000 times
      37: kernel launched 1000 times
          grid: [79]  block: [128]
           device time(us): total=2,006 max=4 min=2 avg=2
           elapsed time(us): total=17,550 max=31 min=15 avg=17
  37: data region reached 2000 times
      37: data copyin transfers: 1000
           device time(us): total=13,797 max=17 min=5 avg=13
      42: data copyout transfers: 1000
           device time(us): total=13,911 max=20 min=5 avg=13
```

The grid is the number of gangs. The block is the size of one gang ($\lceil vector_{length} \times nb_{workers} \rceil$).

Introduction

PGI_ACC_TIME=1 — parallel-data-single.f90

```fortran
28      !$ACC parallel copyout( a(:s) )
29      !$ACC loop
30      do i=1,s
31        a(i) = i
32      enddo
33      !$ACC end parallel
34
35      !$ACC data copy( a(:s) )
36      do j=1,p
37        !$ACC parallel
38        !$ACC loop
39        do i=1,s
40          a(i) = a(i) + 1
41        enddo
42        !$ACC end parallel
43      enddo
44      !$ACC end data
```

```
para NVIDIA devicenum=0
  time(us): 2,069
  28: compute region reached 1 time
    28: kernel launched 1 time
        grid: [79]  block: [128]
        device time(us): total=4 max=4 min=4 avg=4
        elapsed time(us): total=378 max=378 min=378 avg=378
  28: data region reached 2 times
    33: data copyout transfers: 1
        device time(us): total=22 max=22 min=22 avg=22
  35: data region reached 2 times
    35: data copyin transfers: 1
        device time(us): total=14 max=14 min=14 avg=14
    44: data copyout transfers: 1
        device time(us): total=9 max=9 min=9 avg=9
  37: compute region reached 1000 times
    37: kernel launched 1000 times
        grid: [79]  block: [128]
        device time(us): total=2,020 max=5 min=2 avg=2
        elapsed time(us): total=16,670 max=33 min=12 avg=16
```

The grid is the number of gangs. The block is the size of one gang ($\left\lceil vector_{length} \times nb_{workers} \right\rceil$).

Introduction

nvprof or pgprof

- Nvidia Toolkit provides the command line profiler: nvprof;
- **PGI_ACC_TIME=0**, the env. var. is incompatible with nvprof;
- Options:
    - **--cpu-profiling on** : Activates CPU profiling;
    - **--metrics flop_count_dp** : Number of operations;
    - **--metrics dram_read_throughput** : memory read throughput for each kernel running on the CPU;
    - **--metrics dram_write_throughput** : memory write throughput for each kernel running on the CPU;

Introduction

nvprof or pgprof - parallel-data-multi.f90

```
            Type  Time(%)      Time  Calls      Avg      Min      Max  Name
 GPU activities: 47.59%  6.0764ms   1000  6.0760us  4.8320us  24.352us  [CUDA memcpy HtoD]
                 40.21%  5.1337ms   1001  5.1280us  3.8080us  19.935us  [CUDA memcpy DtoH]
                 12.18%  1.5552ms   1000  1.5550us  1.4070us  13.503us  para_37_gpu
     API calls: 48.50%  115.20ms      1  115.20ms  115.20ms  115.20ms  cuDevicePrimaryCtxRetain
                 21.14%  50.208ms      1  50.208ms  50.208ms  50.208ms  cuDevicePrimaryCtxRelease
                 13.10%  31.106ms   3002  10.361us     859ns  61.105us  cuStreamSynchronize
                  5.98%  14.198ms      1  14.198ms  14.198ms  14.198ms  cuMemHostAlloc
                  4.48%  10.639ms      1  10.639ms  10.639ms  10.639ms  cuMemFreeHost
                  2.53%  6.0090ms   1001  6.0030us  4.8850us  28.241us  cuLaunchKernel
                  1.56%  3.6943ms   1000  3.6940us  3.0770us  28.117us  cuMemcpyHtoDAsync
                  1.50%  3.5689ms   1001  3.5650us  2.9210us  23.222us  cuMemcpyDtoHAsync
 OpenACC (excl): 27.17%  21.357ms   2000  10.678us  1.4670us  62.106us  acc_wait@parallel-data-multi.f90:37
                 19.31%  15.178ms   1000  15.177us  4.5620us  40.211us  acc_wait@parallel-data-multi.f90:42
                 18.10%  14.231ms      1  14.231ms  14.231ms  14.231ms  acc_exit_data@parallel-data-multi.f90:29
                  9.42%  7.4051ms   1000  7.4050us  5.7670us  390.34us  acc_enqueue_launch@parallel-data-multi.f90:37 (para_37_gpu)
                  7.67%  6.0306ms   1000  6.0300us  4.6600us  350.63us  acc_enqueue_download@parallel-data-multi.f90:42
                  6.70%  5.2692ms   1000  5.2690us  4.0230us  28.372us  acc_enter_data@parallel-data-multi.f90:37
                  5.71%  4.4901ms   1000  4.4900us  3.7610us  28.943us  acc_enqueue_upload@parallel-data-multi.f90:37
                  3.76%  2.9595ms   1000  2.9590us  2.3790us  21.317us  acc_exit_data@parallel-data-multi.f90:37
                  1.80%  1.4138ms   1000  1.4130us  1.1650us  23.892us  acc_compute_construct@parallel-data-multi.f90:37
```

Introduction

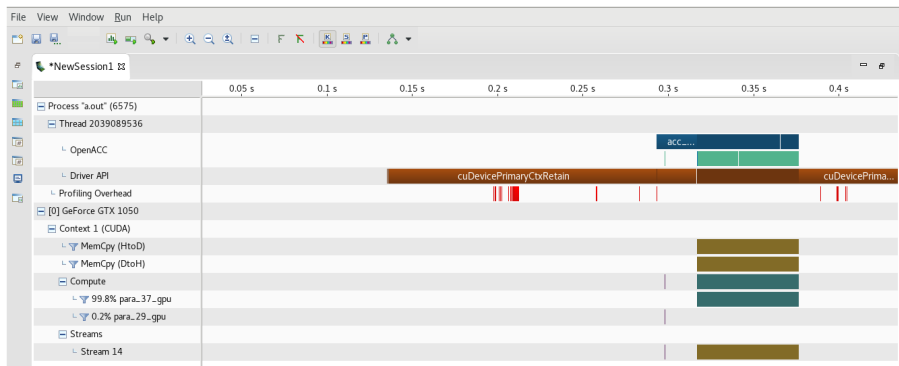## nvprof or pgprof - parallel-data-single.f90

```
            Type  Time(%)      Time  Calls      Avg       Min       Max  Name
GPU activities:   98.72%   1.4593ms   1000  1.4590us  1.4070us  12.447us  para_37_gpu
                   0.69%  10.272us       2  5.1360us  4.8960us  5.3760us  [CUDA memcpy DtoH]
                   0.37%  5.5040us       1  5.5040us  5.5040us  5.5040us  [CUDA memcpy HtoD]
                   0.21%  3.0720us       1  3.0720us  3.0720us  3.0720us  para_28_gpu
     API calls:   64.76%   143.60ms       1  143.60ms  143.60ms  143.60ms  cuDevicePrimaryCtxRetain
                  19.97%   44.288ms       1  44.288ms  44.288ms  44.288ms  cuDevicePrimaryCtxRelease
                   6.45%   14.306ms       1  14.306ms  14.306ms  14.306ms  cuMemHostAlloc
                   3.97%   8.8095ms       1  8.8095ms  8.8095ms  8.8095ms  cuMemFreeHost
                   2.94%   6.5292ms    1004  6.5030us     884ns  27.722us  cuStreamSynchronize
                   1.54%   3.4067ms    1001  3.4030us  3.0620us  23.182us  cuLaunchKernel
 OpenACC (excl):  53.03%   14.340ms       1  14.340ms  14.340ms  14.340ms  acc_exit_data@parallel-data-single.f90:28
                  26.16%   7.0745ms    1000  7.0740us  1.4780us  28.318us  acc_wait@parallel-data-single.f90:37
                  15.58%   4.2143ms    1000  4.2140us  3.7990us  21.628us  acc_enqueue_launch@parallel-data-single.f90:37 (para_37_gpu)
                   3.97%   1.0743ms    1000  1.0740us  1.0140us  16.360us  acc_compute_construct@parallel-data-single.f90:37
```
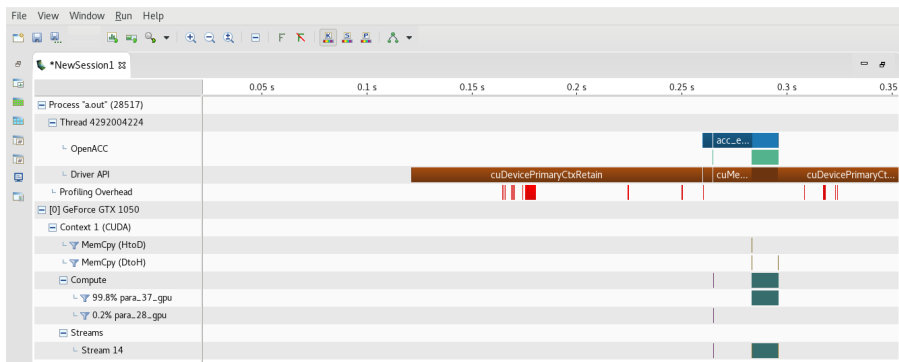
Introduction

pgprof GUI - parallel-data-multi.f90

Introduction

pgprof GUI - parallel-data-single.f90

Parallel Directives

- OpenACC features are activated through directives.
- Directives are treated as comments if the right compiler options are not set.
- The syntax is different for a Fortran or a C/C++ code.

Fortran

```
!$ACC directive <clauses>
...
!$ACC end directive
```

C/C++

```
#pragma acc directive <clauses>
{
    ...
}
```

Directive examples:
kernels, loop, parallel, data, enter data, exit data

Compute Construct

Serial

- The code runs on a single thread;
- 1 gang with 1 worker of size 1.

Kernels

- The compiler analyzes the code and decides the parallelism level of the kernel;
- One kernel is generated for each parallel parallel loop enclosed in the region.

Parallel

- Parallel region to be run on the device;
- Only one kernel is generated;
- The programmer is responsible for sharing the work manually;
- Execution is redundant by default:
    - Gang-redundant;
    - Worker-single;
    - Vector-single.

Serial Directive

Characteristics

- Inform the compiler that the enclosed region has to be offloaded to the GPU;
- Instructions are executed only by one thread, *i.e.* one gang with a worker of size one;
- **Serial** directive is equivalent to the following parameters: **num_gang(1), num_worker(1), vector_length(1)**.

Default Behavior

- Arrays present in the **serial** region not specified in a data clause (**present, copyin, copyout**, etc) or a **declare** directive are assigned to a **copy**. They are **SHARED**;
- Scalar variables are implicitly assigned to a **firstprivate** clause. The are **PRIVATE**.

Serial Directive

Game of Life - Serial – gol.f90

```
23    !$ACC serial
24    do g=1,generations
25      do r=1, rows
26        do c=1, cols
27          old_world(r,c) = world(r,c)
28        enddo
29      enddo
30      do r=1, rows
31        do c=1, cols
32          neigh = old_world(r-1,c-1)+old_world(r,c-1)+&
33          old_world(r+1,c-1)+old_world(r-1,c)+old_world(r+1,c)+&
34          old_world(r-1,c+1)+old_world(r,c+1)+old_world(r+1,c+1)
35          if (old_world(r,c) == 1 .and. (neigh<2.or.neigh>3) ) then
36            world(r,c) = 0
37          else if (neigh == 3) then
38            world(r,c) = 1
39          endif
40        enddo
41      enddo
42      cells = 0
43      do r=1, rows
44        do c=1, cols
45          cells = cells + world(r,c)
46        enddo
47      enddo
48      print *, "Cells_alive_at_generation_", g, ":_", cells
49    enddo
50    !$ACC end serial
```

```
gol:
23, Accelerator serial kernel generated
  Generating Tesla code
24, !$acc do seq
25, !$acc do seq
26, !$acc do seq
30, !$acc do seq
31, !$acc do seq
43, !$acc do seq
44, !$acc do seq
23, Generating implicit copy(world(:,:),old_world(:,:)) [if not already present]
```

Test

- Size : 1000x1000;
- Generations: 100;
- Elapsed time: 123.640s.

Kernels Directive

Characteristics

- Inform the compiler that the enclosed region contains instructions to be offloaded on the device;
- Each loop nest is treated as an independant kernel with its own parameters *i.e.* number of gangs, workers and vector size.

Default Behavior

- Data arrays within **kernels** region that are not specified in a data clause (**present, copyin, copyout**, etc) or a **declare** directive are assigned to a **copy**. They are **SHARED**;
- Scalar variables are implicitly assigned to a **copy** clause. The are **SHARED**.

Important Notes

- Parameters of parallel regions, such as gangs, workers, and vector length, are independent.
- Loop nests are executed consecutively.

Kernels Directive

Game of Life - Kernels – gol.f90

```
23    !$ACC kernels
24    do g=1,generations
25     do r=1, rows
26      do c=1, cols
27        old_world(r,c) = world(r,c)
28      enddo
29     enddo
30     do r=1, rows
31      do c=1, cols
32        neigh = old_world(r-1c-1)+old_world(r,c-1)+&
33        old_world(r+1,c-1)+old_world(r-1c)+old_world(r+1,c)+&
34        old_world(r-1c+1)+old_world(r,c+1)+old_world(r+1,c+1)
35        if (old_world(r,c) == 1 .and. (neigh<2or.neigh>3) ) then
36          world(r,c) = 0
37        else if (neigh == 3) then
38          world(r,c) = 1
39        endif
40      enddo
41     enddo
42     cells = 0
43     do r=1, rows
44      do c=1, cols
45        cells = cells + world(r,c)
46      enddo
47     enddo
48     print *," Cells_alive_at_generation_", g, ":_", cells
49    enddo
50    !$ACC end kernels
```

```
gol:
23, Generating implicit copy(world(:,:),old_world(:,:)) [if not already present]
24, Loop carried dependence due to exposed use of world(:,:) prevents parallelization
    Parallelization would require privatization of array old_world(i2+1,:)
    Generating Tesla code
    24, !$acc loop seq
    25, !$acc loop vector(128) ! threadidx%x
    26, !$acc loop seq
    30, !$acc loop vector(128) ! threadidx%x
    31, !$acc loop seq
    43, !$acc loop vector(128) ! threadidx%x
    44, !$acc loop seq
    45, Generating implicit reduction(+:cells)
25, Loop is parallelizable
26, Loop is parallelizable
30, Loop is parallelizable
31, Loop is parallelizable
43, Loop is parallelizable
44, Loop is parallelizable
```

Test

- Size : 1000x1000;
- Generations: 100;
- Elapsed time: 1.951s.

Parallel Directive

Characteristics

- It creates a parallel region on the device and generates one or more gangs;
- All gangs execute redundantly the instructions within the parallel region, *i.e.* gang-redundant mode;
- Only parallel loop nests with a loop directive are eligible to have their iterations spread among gangs.

loop.f90

```
8     !$ACC parallel
9     a = 0  !!! Gang redundant
10    !$ACC loop  !!! Work sharing
11    do i=1,10000
12      a(i) = i
13    enddo
14    !$ACC end parallel
```

```
loop:
8, Generating Tesla code
    9, !$acc loop vector(128) ! threadidx%x
    11, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
8, Generating implicit copyout(a(:)) [if not already present]
9, Loop is parallelizable
```

Parallel Directive

Default Behavior

- Data arrays within a **parallel** region that are not specified in a data clause (**present, copyin, copyout**, etc) or a **declare** directive are assigned to a **copy**. They are **SHARED**;
- Scalar variables are implicitly assigned to a **firstprivate** clause. The are **PRIVATE**.

Important Notes

- The number of gangs, workers and vector length are constant inside the parallel region.

Parallel Directive

Game of Life - Paralel -acc=noautopar – gol.f90

```
23    !$ACC parallel
24    do g=1,generations
25      do r=1, rows
26        do c=1, cols
27          old_world(r,c) = world(r,c)
28        enddo
29      enddo
30      do r=1, rows
31        do c=1, cols
32          neigh = old_world(r-1,c-1)+old_world(r,c-1)+&
33            old_world(r+1,c-1)+old_world(r-1,c)+old_world(r+1,c)+&
34            old_world(r-1,c+1)+old_world(r,c+1)+old_world(r+1,c+1)
35          if (old_world(r,c) == 1 .and. (neigh<2 or.neigh>3) ) then
36            world(r,c) = 0
37          else if ( neigh == 3 ) then
38            world(r,c) = 1
39          endif
40        enddo
41      enddo
42      cells = 0
43      do r=1, rows
44        do c=1, cols
45          cells = cells + world(r,c)
46        enddo
47      enddo
48      print *, "Cells_alive_at_generation_", g, ":_", cells
49    enddo
50    !$ACC end parallel
```

```
gol:
23, Generating Tesla code
  24, !$acc loop seq
  25, !$acc loop seq
  26, !$acc loop seq
  30, !$acc loop seq
  31, !$acc loop seq
  43, !$acc loop seq
  44, !$acc loop seq
23, Generating implicit copy(old_world(:,:),world(:,:)) [if not already present]
24, Loop carried dependence due to exposed use of world(:,:) prevents parallelization
    Parallelization would require privatization of array old_world(i2+1,:)
25, Loop is parallelizable
26, Loop is parallelizable
30, Loop is parallelizable
31, Loop is parallelizable
43, Loop is parallelizable
44, Loop is parallelizable
```

Test

- Size : 1000x1000;
- Generations: 100;
- Elapsed time: 120.167s

- The sequential code is executed redundantly by all gangs;
- The compiler option **acc=noautopar** is activated to reproduce the expected behavior of the OpenACC specification.

Parallel Directive

Game of Life - Paralel -acc=autopar – gol.f90

```
23    !$ACC parallel
24    do g=1,generations
25      do r=1, rows
26        do c=1, cols
27          old_world(r,c) = world(r,c)
28        enddo
29      enddo
30      do r=1, rows
31        do c=1, cols
32          neigh = old_world(r-1c-1)+old_world(r,c-1)+&
33          old_world(r+1,c-1)+old_world(r-1c)+old_world(r+1,c)+&
34          old_world(r-1c+1)+old_world(r,c+1)+old_world(r+1,c+1)
35          if (old_world(r,c) == 1 .and. (neigh<2or.neigh>3)) then
36            world(r,c) = 0
37          else if (neigh == 3) then
38            world(r,c) = 1
39          endif
40        enddo
41      enddo
42      cells = 0
43      do r=1, rows
44        do c=1, cols
45          cells = cells + world(r,c)
46        enddo
47      enddo
48      print *, "Cells_alive_at_generation_", g, ":_", cells
49    enddo
50    !$ACC end parallel
```

```
gol:
23, Generating Tesla code
   24, !$acc loop seq
   25, !$acc loop seq
   26, !$acc loop vector(128) ! threadidx%x
   30, !$acc loop seq
   31, !$acc loop vector(128) ! threadidx%x
   43, !$acc loop seq
   44, !$acc loop vector(128) ! threadidx%x
   45, Generating implicit reduction(+:cells)
23, Generating implicit copy(old_world(:,:),world(:,:)) [if not already present]
24, Loop carried dependence due to exposed use of world(:,:) prevents parallelization
   Parallelization would require privatization of array old_world(i2+1,:)
25, Loop is parallelizable
26, Loop is parallelizable
30, Loop is parallelizable
31, Loop is parallelizable
43, Loop is parallelizable
44, Loop is parallelizable
```

Test

- Size : 1000x1000;
- Generations: 100;
- Elapsed time: 5.067s

Parallel Directive

Default Behavior

- It is up to the compiler to decide how many workers are generated and their vector size;
- The number of gangs is set at execution time by the runtime;
- Memory is usually the limiting criterion.

Control Clauses

The programmer can set control parameters for **kernels** and **parallel** clauses:

- **num_gangs**: provide number of gangs;
- **num_workers**: provide number of workers;
- **vector_length**: provide vector length.

Important Notes

- These clauses are mainly useful if the code uses a data structure which is difficult for the compiler to analyze;
- The optimal number of gangs is highly dependent on the architecture. Use **num_gangs** with care.

Parallel Directive

parametres_paral.f90

```
14      !$acc parallel num_gangs(10)&
15      !$acc& num_workers(1)        &
16      !$acc& vector_length(128)
17      print *, "Hello_I_am_a_gang"
18      do i=1,1000
19        a(i) = i
20      enddo
21      !$acc end parallel
```

```
1    Hello I am a gang
2    Hello I am a gang
3    Hello I am a gang
4    Hello I am a gang
5    Hello I am a gang
6    Hello I am a gang
7    Hello I am a gang
8    Hello I am a gang
9    Hello I am a gang
10   Hello I am a gang
```

```
param:
14, Generating Tesla code
    18, !$acc loop gang(10), vector(128) ! blockidx%x threadidx%x
14, Generating implicit copyout(a(:)) [if not already present]
18, Loop is parallelizable
```

Parallel Directive

parametres_mod.f90

```
14      !$acc parallel num_gangs(10)&
15      !$acc& num_workers(1)        &
16      !$acc& vector_length(128)
17      print *, "Hello_I_am_a_gang"
18      !$acc loop
19      do i=1,1000
20        a(i) = i
21      enddo
22      !$acc end parallel
```

```
1    Hello I am a gang
2    Hello I am a gang
3    Hello I am a gang
4    Hello I am a gang
5    Hello I am a gang
6    Hello I am a gang
7    Hello I am a gang
8    Hello I am a gang
9    Hello I am a gang
10   Hello I am a gang
```

```
param:
14, Generating Tesla code
    19, !$acc loop gang(10), vector(128) ! blockidx%x threadidx%x
14, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Loop

- Loops are at the heart of OpenACC parallelism;
- The **loop** directive , is responsible for sharing the work, i.e. the iterations of the associated loop;
- It could also activate another level of parallelism.
- Aautomatic nesting is a critical difference between **OpenACC** and **OpenMP**-**GPU** for which the creation of threads relies on the programmer.

Work Sharing – Loop

Clauses

The use of clauses can indicate different parallelism levels:

- **gang**: The iterations of the subsequent loop are distributed block-wise among gangs. It can be gang-redundant or gang-parallel;
- **worker**: Combined with **gangs**, threads of workers are activated, as worker-single or worker-parallel, and the task is shared among those threads.
- **vector**: Workers activate vectors in a single-instruction multiple-thread (SIMT) to share the task in vector single (VS) or vector parallel (VP) modes;
- **seq**: Iterations are sequentially executed on the device;
- **auto**: The compiler analyses the loop region and decides which options are more suitable to respect dependencies;
- **collapse(#loops)**: Merge tightly nested loops;
- **independent**: Tell the compilers that iterations are independent. Useful for the **kernel** directive;
- **private(variable-list)**: Privatize variables;
- **reduction(operation:variable-list)**: Reduction operation.

Work Sharing – Loop

Game of Life – acc=noautopar – Paralel loop – gol.f90

```
23    !$ACC parallel
24    do g=1,generations
25      !$acc loop
26      do r=1, rows
27        do c=1, cols
28          old_world(r,c) = world(r,c)
29        enddo
30      enddo
31      do r=1, rows
32        do c=1, cols
33          neigh = old_world(r-1c-1)+old_world(r,c-1)+&
34            old_world(r+1,c-1)+old_world(r-1c)+old_world(r+1,c)+&
35            old_world(r-1c+1)+old_world(r,c+1)+old_world(r+1,c+1)
36          if (old_world(r,c) == 1 .and. (neigh<2or.neigh>3) ) then
37            world(r,c) = 0
38          else if (neigh == 3) then
39            world(r,c) = 1
40          endif
41        enddo
42      enddo
43      cells = 0
44      do r=1, rows
45        do c=1, cols
46          cells = cells + world(r,c)
47        enddo
48      enddo
49      print *, "Cells_alive_at_generation_", g, ":_", cells
50    enddo
51    !$ACC end parallel
```

```
gol:
23, Generating Tesla code
    24, !$acc loop seq
    26, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
    27, !$acc loop seq
    31, !$acc loop seq
    32, !$acc loop seq
    44, !$acc loop seq
    45, !$acc loop seq
23, Generating implicit copy(old_world(:,:),world(:,:)) [if not already present]
24, Loop carried dependence due to exposed use of world(:,:) prevents parallelization
    Parallelization would require privatization of array old_world(i2+1,:)
27, Loop is parallelizable
31, Loop is parallelizable
32, Loop is parallelizable
44, Loop is parallelizable
45, Loop is parallelizable
```

Work Sharing – Loop

Sequential – loop_serial.f90

```
38      !$ACC serial
39      do i=1,10000
40        a(i) = i
41      enddo
42      !$ACC end serial
```

- Active threads: 1;
- Number of operations: *nx*.


```
loop :
38, Accelerator serial kernel generated
    Generating Tesla code
    39, !$acc do seq
38, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Gang-Redundant Worker-Serial Vector-Single (GRWSVS) –loop_grwsvs.f90

```
39      !$ACC parallel num_gangs(10)
40      do i=1,10000
41        a(i) = i
42      enddo
43      !$ACC end parallel
```


- Redundant execution by gang leaders;
- Active threads: 10;
- Number of operations: $10 \times nx$.

```
loop:
39, Generating Tesla code
    40, !$acc loop gang(10), vector(128) ! blockidx%x threadidx%x
39, Generating implicit copyout(a(:)) [if not already present]
40, Loop is parallelizable
```

Work Sharing – Loop

Gang-Parallel Worker-Single Vector-Single (GPWSVS) – loop_gpwsvs.f90

```
40      !$ACC parallel num_gangs(10)
41      !$ACC loop gang
42      do i=1,10000
43        a(i) = i
44      enddo
45      !$ACC end parallel
```


- Each gang executes a different block of iterations loop;
- Active threads: 10;
- Number of operations: nx.

```
loop:
40, Generating Tesla code
    42, !$acc loop gang(10), vector(128) ! blockidx%x threadidx%x
40, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Gang-Parallel Worker-Parallel Vector-Single (GPWPVS) – loop_gpwpvs.f90

```
40      !$ACC parallel num_gangs(10)
41      !$ACC loop gang worker
42      do i=1,10000
43        a(i) = i
44      enddo
45      !$ACC end parallel
```


- Iterations are shared among the active workers of each gang;
- Active threads: $10 \times nb_{workers}$;
- Number of operations: nx.

```
loop:
40, Generating Tesla code
    42, !$acc loop gang(10) ! blockidx%x threadidx%y
40, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Gang-Parallel Worker-Single Vector-Parallel (GPWSVP) – loop_gpwsvp.f90

```
40      !$ACC parallel num_gangs(10)
41      !$ACC loop gang vector
42      do i=1,10000
43        a(i) = i
44      enddo
45      !$ACC end parallel
```

- Iterations are shared among the threads of the worker of all gangs;
- Active threads: $10 \times vector_{length}$;
- Number of operations: nx.


```
loop:
40, Generating Tesla code
    42, !$acc loop gang(10), vector(128) ! blockidx%x threadidx%x
40, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Gang-Parallel Worker-Parallel Vector-Parallel (GPWPVP) – loop_gpwpvp.f90

```
42      !$ACC parallel num_gangs(10)
43      !$ACC loop gang worker vector
44      do i=1,10000
45        a(i) = i
46      enddo
47      !$ACC end parallel
```

- Iterations are shared among the threads of the worker of all gangs;
- Active threads: $10 \times nb_{workers} \times vector_{length}$;
- Number of operations: nx.


```
loop:
42, Generating Tesla code
     44, !$acc loop gang(10), worker(4), vector(32) ! blockidx%x threadidx%y threadidx%x
42, Generating implicit copyout(a(:)) [if not already present]
```

Work Sharing – Loop

Gang-Reduntant Worker-Parallel Vector-Single (GRWPVS) – loop_grwpvs.f90

```
44      !$ACC parallel num_gangs(10)
45      !$ACC loop worker
46      do i=1,10000
47        a(i) = i
48      enddo
49      !$ACC end parallel
```


- All iterations are assigned to each active gang which share the task among workers;
- Active threads: $10 \times nb_{workers}$;
- Number of operations: $10 \times nx$.

```
loop:
44, Generating Tesla code
    46, !$acc loop gang(10) ! blockidx%x threadidx%y
44, Generating implicit copyout(a(:)) [if not already present]
46, Loop is parallelizable
```

Work Sharing – Loop

Gang-Reduntant Worker-Single Vector-Parallel (GRWPVS) – loop_grwsvp.f90

```fortran
41       !$ACC parallel num_gangs(10)
42       !$ACC loop vector
43       do i=1,10000
44          a(i) = i
45       enddo
46       !$ACC end parallel
```


- All iterations are assigned to each active gang which share the task among vectors of one worker;
- Active threads: $10 \times vector_{length}$;
- Number of operations: $10 \times nx$.

```
loop:
41, Generating Tesla code
     43, !$acc loop vector(128) ! threadidx%x
41, Generating implicit copyout(a(:)) [if not already present]
43, Loop is parallelizable
```

Work Sharing – Loop

Gang-Reduntant Worker-Parallel Vector-Parallel (GRWPVP) – loop_grwpvp.f90

```
49      !$ACC parallel num_gangs(10)
50      !$ACC loop worker vector
51      do i=1,10000
52        a(i) = i
53      enddo
54      !$ACC end parallel
```


- All iterations are assigned to each active gang which share the task among the whole thread grid;
- Active threads: $10 \times nb_{workers} \times vector_{length}$;
- Number of operations: $10 \times nx$.

```
loop:
49, Generating Tesla code
    51, !$acc loop gang(10), worker(4), vector(32) ! blockidx%x threadidx%y threadidx%x
49, Generating implicit copyout(a(:)) [if not already present]
51, Loop is parallelizable
```

Work Sharing – Loop

Reminder

- There is no thread synchronization at gang level (specially at the end of a loop directive). **Risk of race condition**;
- Synchronization is present when using the *loop* directive with **worker** and/or **vector** parallelism since the threads of a gang wait until the end of the iterations they execute to start a new portion of the code after the loop;
- **Gang** parallelism should be avoided inside parallel regions in the presence of dependencies.

pb-sync.f90

```
!$acc parallel
!$acc loop gang
do i=1,nx
  a(i) = 1.0_8
enddo
!$acc loop gang reduction(+:summ)
do i=nx,1,-1
  summ = summ + a(i)
enddo
!$acc end parallel
```

corr_sync.f90

```
!$acc parallel
!$acc loop worker vector
do i=1,nx
  a(i) = 1.0_8
enddo
!$acc loop worker vector reduction(+:summ)
do i=nx,1,-1
  summ = summ + a(i)
enddo
!$acc end parallel
```

- $Summ = 5.0066080E + 07$ (Race condition);

- $Summ = 1.0000000E + 08$

Merging Directives

Kernels, Parallel and Loop

- One can merge read **Kernels**/ **Parallel** regions with the **loop** directive;
- The clauses available for this construct are those of both constructs.

fused.f90

```
18      !$acc kernels
19      do i=1,10000
20        a1(i) = i
21        a2(i) = i
22        b(i) = i+i*i
23      enddo
24      !$acc end kernels
25      !$acc kernels
26      !$acc loop
27      do i=1,10000
28        a1(i) = b(i)*2
29        summ1 = summ1 + a1(i)
30      enddo
31      !$acc end kernels
32      !$acc kernels loop worker vector
33      do i=1,10000
34        a2(i) = b(i)*2
35        summ2 = summ2 + a2(i)
36      enddo
```

```
fused:
18, Generating implicit copyout(a1(:),a2(:),b(:)) [if not already present]
19, Loop is parallelizable
    Generating Tesla code
    19, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
25, Generating implicit copyin(b(:)) [if not already present]
27, Loop is parallelizable
    Generating Tesla code
    27, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
    29, Generating implicit reduction(+:summ1)
32, Generating implicit copyin(b(:)) [if not already present]
    Generating implicit copyout(a2(:)) [if not already present]
33, Loop is parallelizable
    Generating Tesla code
    33, !$acc loop gang, worker(4), vector(32) ! blockidx%x threadidx%y threadidx%x
    35, Generating implicit reduction(+:summ2)
```

- $summ1 = 1146749120$
- $summ2 = 1146749120$

Reductions

Reduction directive

- A variable assigned to a **reduction** clause is privatised for each element of the parallelism level of the loop;
- At the end of the region, an operation is executed using options provided below:

Reduction operations in Fortran.

Operation	Effect
+	Summ
*	Product
max	Maximum
min	Minimum
iand	Bitwise and
ior	Bitwise or
ieor	Bitwise xor
.and.	Logical and
.or.	Logical or

Reduction operations in C/C++.

Operation	Effect
+	Summ
*	Product
max	Maximum
min	Minimum
&	Bitwise and
|	Bitwise or
&&	Logical and
||	Logical or

Restrictions
The variable have to be a scalar with a numerical value:

- Fortran: **integer**,**real**,**double prestion**,**complex**;
- C: **char**,**int**,**float**,**double**, **_Complex**;
- C++: **char**,**wchar_t**,**int**,**float**,**double**;

In the OpenACC specification 2.7, reductions are possible on arrays but the implementation is lacking in PGI (for the moment).

Reductions

Example - corr_sync.f90

```fortran
18    program reduction
19      real     :: a(100000000)
20      integer :: i
21      integer :: nx=100000000
22      real     :: summ=0
23      !$acc parallel
24      !$acc loop worker vector
25      do i=1,nx
26        a(i) = 1.0_8
27      enddo
28      !$acc loop worker vector reduction(+:summ)
29      do i=nx,1,-1
30        summ = summ + a(i)
31      enddo
32      !$acc end parallel
33      write (*,*) summ
34    end program reduction
```

Memory Access and Coalescing

Principles

- Contiguous access to memory by the threads of a worker can be merged. This optimizes the use of memory bandwidth;
- This happens if thread i reaches memory location n, and thread $i + 1$ reaches memory location $n + 1$ and so on;
- For loop nests, the loop which has **vector** parallelism should have contiguous access to memory.

Contiguous data access.

Non-contiguous data access.

Memory Access and Coalescing

Sources

loop_nocoalescing.f90

```
22      !$ACC parallel
23      !$ACC loop gang
24      do i=1,nx
25        !$ACC loop vector
26        do j=1,nx
27          a(i,j) = 1.14d-8
28        enddo
29      enddo
30      !$ACC end parallel
```

loop_coalescing-seq.f90

```
22      !$ACC parallel
23      !$ACC loop gang vector
24      do i=1,nx
25        !$ACC loop seq
26        do j=1,nx
27          a(i,j) = 1.14d-8
28        enddo
29      enddo
30      !$ACC end parallel
```

$$nx = 10000$$

loop_coalescing-vec.f90

```
22      !$ACC parallel
23      !$ACC loop gang
24      do j=1,nx
25        !$ACC loop vector
26        do i=1,nx
27          a(i,j) = 1.14d-8
28        enddo
29      enddo
30      !$ACC end parallel
```

Memory Access and Coalescing

Compiling

loop_nocoalescing.f90

```
loop:
22, Generating Tesla code
    24, !$acc loop gang ! blockidx%x
    26, !$acc loop vector(128) ! threadidx%x
22, Generating implicit copyout(a(1:nx,1:nx)) [if not already present]
26, Loop is parallelizable
```

loop_coalescing-seq.f90

```
loop:
22, Generating Tesla code
    24, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
    26, !$acc loop seq
22, Generating implicit copyout(a(1:nx,1:nx)) [if not already present]
```

loop_coalescing-vec.f90

```
loop:
22, Generating Tesla code
    24, !$acc loop gang ! blockidx%x
    26, !$acc loop vector(128) ! threadidx%x
22, Generating implicit copyout(a(1:nx,1:nx)) [if not already present]
26, Loop is parallelizable
```

Memory Access and Coalescing

Computing

loop_nocoalescing.f90

```
loop  NVIDIA  devicenum=0
time(us): 126,451
22: compute region reached 1 time
    22: kernel launched 1 time
        grid: [10000]  block: [128]
         device time(us): total=94,509 max=94,509 min=94,509 avg=94,509
         elapsed time(us): total=94,553 max=94,553 min=94,553 avg=94,553
22: data region reached 2 times
    30: data copyout transfers: 24
         device time(us): total=31,942 max=1,394 min=1,211 avg=1,330
```

loop_coalescing-seq.f90

```
loop  NVIDIA  devicenum=0
time(us): 37,550
22: compute region reached 1 time
    22: kernel launched 1 time
        grid: [79]  block: [128]
         device time(us): total=5,634 max=5,634 min=5,634 avg=5,634
         elapsed time(us): total=5,677 max=5,677 min=5,677 avg=5,677
22: data region reached 2 times
    30: data copyout transfers: 24
         device time(us): total=31,916 max=1,505 min=1,174 avg=1,329
```

loop_coalescing-vec.f90

```
loop  NVIDIA  devicenum=0
time(us): 35,804
22: compute region reached 1 time
    22: kernel launched 1 time
        grid: [10000]  block: [128]
         device time(us): total=4,145 max=4,145 min=4,145 avg=4,145
         elapsed time(us): total=4,188 max=4,188 min=4,188 avg=4,188
22: data region reached 2 times
    30: data copyout transfers: 24
         device time(us): total=31,659 max=1,337 min=1,126 avg=1,319
```

	W/O. Coal.	With Coal.
Time (μs)	≈ 126	≈ 36

61 / 89

Memory Access and Coalescing

Computing

loop_nocoalescing.f90 − gang(i)/vector(j)

loop_coalescing-seq.f90 − gang-vector(i)/seq(j)

loop_coalescing-vec.f90 − gang(j)/vector(i)

Memory Access and Coalescing

Computing

loop_nocoalescing.f90 – gang(i)/vector(j)

loop_coalescing-seq.f90 – gang-vector(i)/seq(j)

loop_coalescing-vec.f90 – gang(j)/vector(i)

Memory Access and Coalescing

Computing

loop_nocoalescing.f90 – gang(i)/vector(j)

**loop_coalescing-seq.f90 – gang-vector(i)/seq(j)**

loop_coalescing-vec.f90 – gang(j)/vector(i)

Memory Access and Coalescing

Computing

loop_nocoalescing.f90 – gang(i)/vector(j)

loop_coalescing-seq.f90 – gang-vector(i)/seq(j)

**loop_coalescing-vec.f90 – gang(j)/vector(i)**

Memory Access and Coalescing

Computing

loop_nocoalescing.f90 – gang(i)/vector(j)

loop_coalescing-seq.f90 – gang-vector(i)/seq(j)

loop_coalescing-vec.f90 – gang(j)/vector(i)

Routines

- Routines and functions need to be declared using the **routine** directive;
- The **routine** directive provides information to the compiler that a device version of the function/subroutine has to be generated;
- It is mandatory to set the parallelism level inside the function (**seq, gang, worker,vector**).

routine-wrong.f90

```
22      !$ACC parallel
23      !$ACC loop
24      do i=1,s
25        call fill( a(:,:), s, i )
26      enddo
27      !$ACC end parallel
28      write(*,*) a(1,10)
29      contains
30      subroutine fill( arr, j, k )
31        integer, intent(out) :: arr(:,:)
32        integer, intent(in)  :: j, k
33        integer :: l
34        do l=1,j
35          arr(k,l) = 2
36        enddo
37      end subroutine
```

```
PGF90-S-0155-Procedures called in a compute region must have &
            acc routine information: fill (routine-wrong.f90: 25)
PGF90-S-0155-Accelerator region ignored;                      &
            see -Minfo messages (routine-wrong.f90: 22)
routine:
22, Accelerator region ignored
25, Accelerator restriction: call to 'fill' with no acc routine information
0 inform,   0 warnings,   2 severes, 0 fatal for routine
```

Routines

- Routines and functions need to be declared using the **routine** directive;
- The **routine** directive provides information to the compiler that a device version of the function/subroutine has to be generated;
- It is mandatory to set the parallelism level inside the function (**seq, gang, worker,vector**).

routine-corr.f90

```
22        !$ACC parallel copyout(a)
23        !$ACC loop
24        do i=1,s
25          call fill( a(:,:), s, i )
26        enddo
27        !$ACC end parallel
28        write(*,*) a(1,10)
29        contains
30        subroutine fill( arr, j, k )
31          !$ACC routine seq
32          integer, intent(out) :: arr(:,:)
33          integer, intent(in) :: j, k
34          integer :: l
35          do l=1,j
36            arr(k,l) = 2
37          enddo
38        end subroutine
```

```
routine:
22. Generating copyout(a(:,:)) [if not already present]
    Generating implicit copy(.S0000) [if not already present]
    Generating Tesla code
    24, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
fill:
30. Generating acc routine seq
    Generating Tesla code
```

Data on Device

Opening a Data Region

- There are several ways of making data visible on devices by opening different kinds of data regions.

Computation offloading
Routines and functions need to be declared using the

- **serial**;
- **parallel**;
- **kernel**.

Global region
An implicit data region is opened during the lifetime of the program. The management of this region is done with the use of the **enter data** and **exit data** directives.

Local regions
To open a data region inside a programming unit (function, subroutine) it is necessary to use the **data** directive inside a code block.

Data region associated to programming unit lifetime
A data region is created when a procedure is called (function or subroutine). It is available during the lifetime of the procedure. To make data visible use the **declare** directive.

Notes
The actions taken for the data inside these regions depend on the clause in which they appear.

Data on Device

Data Clauses

Abbreviations and definitions:

- H : Host;
- D : Device;

- Variable: The variable can be a scalar or an array as well.

Data Movement

- **copyin** : The variable is copied $H \rightarrow D$. Memory is allocated when entering the the region;
- **copyout** : The variable is copied $D \rightarrow H$. Memory is allocated when entering the the region;
- **copy** : **copyin** + **copyout**.

No Data Movement

- **create** : The memory allocated when entering the region;
- **present** : The variable is already in the device;
- **delete** : It dellocates memory on the device used by the variable.

Note

By default, the clauses check if the variable is already on the device. If so, no action is taken. It is possible to see clauses prefixed with **present_or_** or **p** for OpenACC 2.0 compatibility.

Other Clauses from the OpenACC 2.7 Standard.

- **deviceptr (Sec. 2.7.3, pg. 41, OpenACC-2.7 Std.);**
- **no_create (Sec. 2.7.9, pg. 44, OpenACC-2.7 Std.);**

- **attach (Sec. 2.7.11, pg. 45, OpenACC-2.7 Std.);**
- **dettach (Sec. 2.7.12, pg. 45, OpenACC-2.7 Std.).**

Data on Device

Shape of Arrays

- It is necessary to specify the shape of an array when transferring data;
- Fortran and C++ do not use the same syntax when transferring arrays.

Fortran − array-shape.f90

```
22      ! Copy a 2−D array on the GPU − matrix "a"
23      !$ACC parallel loop gang copy(a(1:s,1:s))
24      do j=1,s
25        !$ACC loop worker vector
26        do i=1,s
27          a(i,j) = 0
28        enddo
29      enddo
30      ! Copyout columns 100 to 199 included
31      ! to the host
32      !$ACC parallel loop gang copy(a(1:s,100:199))
33      do j=100,199
34        !$ACC loop worker vector
35        do i=1,s
36          a(i,j) = 42
37        enddo
38      enddo
```

- The array shape have to be provided in parentheses;
- It is necessary to provide the first and last indices.

Data on Device

Shape of Arrays

- It is necessary to specify the shape of an array when transferring data;
- Fortran and C++ do not use the same syntax when transferring arrays.

Fortran – array-shape.f90

```
array_shape:
23, Generating copy(a(1:s,1:s)) [if not already present]
     Generating Tesla code
     24, !$acc loop gang ! blockidx%x
     26, !$acc loop worker(4), vector(32) ! threadidx%y threadidx%x
26, Loop is parallelizable
32, Generating copy(a(1:s,100:199)) [if not already present]
     Generating Tesla code
     33, !$acc loop gang ! blockidx%x
     35, !$acc loop worker(4), vector(32) ! threadidx%y threadidx%x
35, Loop is parallelizable
```

- The array shape have to be provided in parentheses;
- It is necessary to provide the first and last indices.

Data on Device

Shape of Arrays

- It is necessary to specify the shape of an array when transferring data;
- Fortran and C++ do not use the same syntax when transferring arrays.

C/C++ – array-shape.cpp

```
28    // Copy the array "a" by giving first element
29    // and the size of the array
30    #pragma acc parallel loop gang copy (a[0:s][0:s])
31    for (int i=0; i<s; ++i )
32      #pragma acc loop worker vector
33      for ( int j=0; j<s; ++j )
34        a[i][j]=0;
35    // Copy copy columns 99 to 198
36    #pragma acc parallel loop gang copy (a[0:s][99:100])
37    for (int i=0; i<s; ++i)
38      #pragma acc loop worker vector
39      for (int j=99; j<199; ++j)
40        a[i][j]=42;
```

- The array shape have to be provided in square brackets;
- It is necessary to provide the first index and the number of elements.

Data on Device

Shape of Arrays

- It is necessary to specify the shape of an array when transferring data;
- Fortran and C++ do not use the same syntax when transferring arrays.

C/C++ − array-shape.cpp

```
main :
26 , Generating copy(a[:s][:s]) [if not already present]
     Generating Tesla code
     31 , #pragma acc loop gang /* blockIdx.x */
     33 , #pragma acc loop worker(4), vector(32) /* threadIdx.y threadIdx.
33 , Loop is parallelizable
34 , Generating copy(a[:s][99:100]) [if not already present]
     Generating Tesla code
     37 , #pragma acc loop gang /* blockIdx.x */
     39 , #pragma acc loop worker(4), vector(32) /* threadIdx.y threadIdx.
39 , Loop is parallelizable
```

- The array shape have to be provided in square brackets;
- It is necessary to provide the first index and the number of elements.

Data on Device

Restrictions

- In Fortran, the last index of an assumed-size dummy array must be specified;
  The dummy argument is a deferred-shape array with (:) bounds.
- In C/C++, the number of elements of a dynamically allocated array must be specified.

Notes

- The shape must be specified when using a slice;
  A slice is subset of elements from an array which is rearranged into another array.
- If the first index is omitted, it is considered as the default of the language
  - **Fortran:** 1;
  - **C/C++:** 0.

Data Region

Parallel Regions

Compute constructs **serial**, **parallel**, **kernels** have a data region associated with variables which are necessary to execution.

parallel-data.f90

```
20      !$ACC parallel copyout( a(:10000) )
21      !$ACC loop
22      do i=1,10000
23        a(i) = i
24      enddo
25      !$ACC end parallel
```

```
parallel_data:
20, Generating copyout(a(:)) [if not already present]
    Generating Tesla code
22, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
```

```
parallel_data  NVIDIA  devicenum=0
time(us): 30
20: compute region reached 1 time
 20: kernel launched 1 time
  grid: [79]  block: [128]
  device time(us): total=4 max=4 min=4 avg=4
  elapsed time(us): total=373 max=373 min=373 avg=373
20: data region reached 2 times
 25: data copyout transfers: 1
  device time(us): total=26 max=26 min=26 avg=26
```

```
                Type  Time(%)  Name
GPU activities:  67.24%  [CUDA memcpy DtoH]
                 32.76%  parallel_data_20_gpu
   API calls:  65.14%  cuDevicePrimaryCtxRetain
               21.76%  cuDevicePrimaryCtxRelease
                7.61%  cuMemHostAlloc
                5.08%  cuMemFreeHost
```

Data region
Parallel region

Data Region

parallel-data-multi.f90 – p=1000 and s=10000

```
28      !$ACC parallel copyout( a(:s) )
29      !$ACC loop
30      do i=1,s
31        a(i) = i
32      enddo
33      !$ACC end parallel
34      do j=1,p
35        !$ACC parallel copy( a(:s) )
36        !$ACC loop
37        do i=1,s
38          a(i) = a(i) + 1
39        enddo
40        !$ACC end parallel
41      enddo
```

```
time(us): 26,006
28: compute region reached 1 time
  28: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=3 max=3 min=3 avg=3
    elapsed time(us): total=372 max=372 min=372 avg=372
28: data region reached 2 times
  33: data copyout transfers: 1
    device time(us): total=21 max=21 min=21 avg=21
35: compute region reached 1000 times
  35: kernel launched 1000 times
    grid: [79]  block: [128]
    device time(us): total=2,032 max=12 min=2 avg=2
    elapsed time(us): total=19,385 max=56 min=15 avg=19
35: data region reached 2000 times
  35: data copyin transfers: 1000
    device time(us): total=11,589 max=36 min=7 avg=11
  40: data copyout transfers: 1000
    device time(us): total=12,361 max=40 min=7 avg=12
```

```
parallel_data:
28, Generating copyout(a(:s)) [if not already present]
    Generating Tesla code
    30, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
35, Generating copy(a(:s)) [if not already present]
    Generating Tesla code
    37, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
```

Notes

- Compute region reached 1000 times;
- Data region reached 2000 times (**copyin**+**copyout**).

Data Region

parallel-data-single.f90 – p=1000 and s=10000

```
27      !$ACC parallel copyout( a(:s) )
28      !$ACC loop
29      do i=1,s
30        a(i) = i
31      enddo
32      !$ACC end parallel
33      !$ACC data copy( a(:s) )
34      do j=1,p
35        !$ACC parallel
36        !$ACC loop
37        do i=1,s
38          a(i) = a(i) + 1
39        enddo
40        !$ACC end parallel
41      enddo
42      !$ACC end data
```

```
parallel_data :
27, Generating copyout(a(:s)) [if not already present]
    Generating Tesla code
    29, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
33, Generating copy(a(:s)) [if not already present]
35, Generating Tesla code
    37, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
```

```
time(us): 2,076
27: compute region reached 1 time
  27: kernel launched 1 time
    grid: [79]  block: [128]
     device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=357 max=357 min=357 avg=357
27: data region reached 2 times
  32: data copyout transfers: 1
     device time(us): total=20 max=20 min=20 avg=20
33: data region reached 2 times
  33: data copyin transfers: 1
     device time(us): total=15 max=15 min=15 avg=15
  42: data copyout transfers: 1
     device time(us): total=9 max=9 min=9 avg=9
35: compute region reached 1000 times
  35: kernel launched 1000 times
    grid: [79]  block: [128]
     device time(us): total=2,028 max=12 min=2 avg=2
    elapsed time(us): total=17,143 max=47 min=13 avg=17
```

Notes

- Compute region reached 1000 times;
- Data region reached 3 times.

Data Region

optm.f90 – p=1000 and s=10000

```
27      !$ACC data copyout( a(:s) )
28      !$ACC parallel loop
29      do i=1,s
30        a(i) = i
31      enddo
32      do j=1,p
33        !$ACC parallel loop
34        do i=1,s
35          a(i) = a(i) + 1
36        enddo
37      enddo
38      !$ACC end data
```

```
parallel_data:
28, Generating copyout(a(:s)) [if not already present]
29, Generating Tesla code
30, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
34, Generating Tesla code
35, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
```

```
time(us): 2,040
28: data region reached 2 times
  39: data copyout transfers: 1
    device time(us): total=23 max=23 min=23 avg=23
29: compute region reached 1 time
  29: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=379 max=379 min=379 avg=379
34: compute region reached 1000 times
  34: kernel launched 1000 times
    grid: [79]  block: [128]
    device time(us): total=2,013 max=5 min=2 avg=2
    elapsed time(us): total=16,992 max=29 min=14 avg=16
```

Notes

- Compute region reached 1000 times;
- Data region reached 2 times.

Data Region

Visual Example – Multiple data and parallel regions

- 8 transfers;
- 2 allocations.

Data Region

Visual Example – Multiple parallel regions within a single data region

- 4 transfers;
- 1 allocations;
- A, B and C are now transferred at entry and exit of the data region.

Data Region

Visual Example – Using **update** in parallel and data regions

- 6 transfers and 1 allocation;
- Clauses check for data presence. However, it is a good practice to use the **present** clause in order to make the code clear;
- The **update** clause can be used to make sure data is up to date in the host or device.
    - The **update** clause is used inside a data region;
    - The **update** clause cannot be used inside a parallel region.

Data Region

Update – **self** or **host** – update-err.f90 – p=42, s=1000

```
29    !$ACC data copyout( a(:s) )
30    !$ACC parallel loop
31    do i=1,s
32      a(i) = 0
33    enddo
34    do j=1,p
35      call random_number(test)
36      rng = floor(test*100)
37      !$ACC parallel loop copyin(rng) &
38      !$ACC& copyout(a)
39      do i=1,s
40        a(i) = a(i) + rng
41      enddo
42    enddo
43    ! write(*,*) "before update self", a(p)
44    !!$ACC update self(a(p:p))
45    ! write(*,*) "after update self", a(p)
46    !$ACC serial
47    a(p) = p
48    !$ACC end serial
49    write(*,*) "before_end_data", a(p)
50    !$ACC end data
51    write(*,*) "after_end_data", a(p)
```

The **self** and **host** clauses update the variable in the
$H \rightarrow D$ direction.

```
before end data          0
after end data          42
```

```
para:
29, Generating copyout(a(:s)) [if not already present]
30, Generating Tesla code
   31, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
37, Generating copyout(a(:)) [if not already present]
   Generating copyin(rng) [if not already present]
   Generating Tesla code
   39, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
46, Accelerator serial kernel generated
   Generating Tesla code
```

The **"a"** array is not initialized on the host
before the end of the data region in
update-err.f90.

75 / 89

Data Region

Update – **self** or **host** – update-err.f90 – p=42, s=1000

```
29    !$ACC data copyout( a (:s) )
30    !$ACC parallel loop
31    do i =1,s
32      a ( i ) = 0
33    enddo
34    do j =1,p
35      call random_number(test)
36      rng = floor(test *100)
37      !$ACC parallel loop copyin(rng) &
38      !$ACC& copyout(a)
39      do i =1,s
40        a ( i ) = a ( i ) + rng
41      enddo
42    enddo
43    ! write(*,*) "before update self", a(p)
44    !!$ACC update self(a(p:p))
45    ! write(*,*) "after update self", a(p)
46    !$ACC serial
47    a(p) = p
48    !$ACC end serial
49    write(*,*) "before_end_data", a(p)
50    !$ACC end data
51    write(*,*) "after_end_data", a(p)
```

The **self** and **host** clauses update the variable in the $H \rightarrow D$ direction.

```
before end data          0
after end data          42
```

```
para NVIDIA devicenum=0
time(us): 497
29: data region reached 2 times
  50: data copyout transfers: 1
   device time(us): total=22 max=22 min=22 avg=22
30: compute region reached 1 time
  30: kernel launched 1 time
   grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=362 max=362 min=362 avg=362
37: compute region reached 42 times
  37: kernel launched 42 times
   grid: [79]  block: [128]
    device time(us): total=86 max=3 min=2 avg=2
    elapsed time(us): total=738 max=27 min=15 avg=17
37: data region reached 84 times
  37: data copyin transfers: 42
   device time(us): total=383 max=10 min=8 avg=9
46: compute region reached 1 time
  46: kernel launched 1 time
   grid: [1]  block: [1]
    device time(us): total=2 max=2 min=2 avg=2
    elapsed time(us): total=17 max=17 min=17 avg=17
```

The **"a"** array is not initialized on the host before the end of the data region in update-err.f90.

Data Region

Update – **self** or **host** – update-corr.f90 – p=42, s=1000

```
28    !$ACC data copyout( a(:s) )
29    !$ACC parallel loop
30    do i=1,s
31      a(i) = 0
32    enddo
33    do j=1,p
34      call random_number(test)
35      rng = floor(test*100)
36      !$ACC parallel loop copyin(rng) &
37      !$ACC& copyout(a)
38      do i=1,s
39        a(i) = a(i) + rng
40      enddo
41    enddo
42    write(*,*) "before_update_self", a(p)
43    !$ACC update self(a(p:p))
44    write(*,*) "after_update_self", a(p)
45    !$ACC serial
46    a(p) = p
47    !$ACC end serial
48    write(*,*) "before_end_data", a(p)
49    !$ACC update host(a(p:p))
50    write(*,*) "second_update_host", a(p)
51    !$ACC end data
52    write(*,*) "after_end_data", a(p)
```

The **self** and **host** clauses update the variable in the $H \rightarrow D$ direction.

```
before update self            0
after update self          2259
before end data            2259
second update host           42
after end data               42
```

```
para:
    28, Generating copyout(a(:s)) [if not already present]
    29, Generating Tesla code
        30, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
    36, Generating copyout(a(:)) [if not already present]
        Generating copyin(rng) [if not already present]
        Generating Tesla code
        38, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
    43, Generating update self(a(p))
    45, Accelerator serial kernel generated
        Generating Tesla code
    49, Generating update self(a(p))
```

The **"a"** array is initialized on the host after the **update** directive.

Data Region

Update – **self** or **host** – update-corr.f90 – p=42, s=1000

```
28    !$ACC data copyout( a(:s) )
29    !$ACC parallel loop
30    do i=1,s
31      a(i) = 0
32    enddo
33    do j=1,p
34      call random_number(test)
35      rng = floor(test*100)
36      !$ACC parallel loop copyin(rng) &
37      !$ACC& copyout(a)
38      do i=1,s
39        a(i) = a(i) + rng
40      enddo
41    enddo
42    write(*,*) "before_update_self", a(p)
43    !$ACC update self(a(p:p))
44    write(*,*) "after_update_self", a(p)
45    !$ACC serial
46    a(p) = p
47    !$ACC end serial
48    write(*,*) "before_end_data", a(p)
49    !$ACC update host(a(p:p))
50    write(*,*) "second_update_host", a(p)
51    !$ACC end data
52    write(*,*) "after_end_data", a(p)
```

The **self** and **host** clauses update the variable in the
$H \rightarrow D$ direction.

```
before update self           0
after update self         2259
before end data           2259
second update host          42
after end data              42
```

```
para  NVIDIA devicenum=0
time(us): 517
28: data region reached 2 times
  51: data copyout transfers: 1
      device time(us): total=9 max=9 min=9 avg=9
29: compute region reached 1 time
  29: kernel launched 1 time
      grid: [79]  block: [128]
      device time(us): total=3 max=3 min=3 avg=3
      elapsed time(us): total=407 max=407 min=407 avg=407
36: compute region reached 42 times
  36: kernel launched 42 times
      grid: [79]  block: [128]
      device time(us): total=86 max=3 min=2 avg=2
      elapsed time(us): total=745 max=38 min=16 avg=17
36: data region reached 84 times
  36: data copyin transfers: 42
      device time(us): total=391 max=18 min=8 avg=9
43: update directive reached 1 time
  43: data copyout transfers: 1
      device time(us): total=21 max=21 min=21 avg=21
45: compute region reached 1 time
  45: kernel launched 1 time
      grid: [1]  block: [1]
      device time(us): total=3 max=3 min=3 avg=3
      elapsed time(us): total=32 max=32 min=32 avg=32
49: update directive reached 1 time
  49: data copyout transfers: 1
      device time(us): total=4 max=4 min=4 avg=4
```

The **"a"** array is initialized on the host after the **update**
directive.

Data Region

Update – **device**

The **device** clause updates the variable in the $D \to H$ direction.

Data Region

Global Data Regions – **declare**

Important Notes The lifetime of data inside the **declare** directive is the same lifetime of the scope of the code region where it is used. Example:

Zone	Scope
Module	Summ
Function	Product
Subroutine	Maximum

```
loop:
13, Generating Tesla code
    14, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
13, Generating implicit copyout(a(:)) [if not already present]
```

```
Accelerator Kernel Timing data
loop  NVIDIA  devicenum=0
time(us): 15
 13: compute region reached 1 time
  13: kernel launched 1 time
    grid: [79]  block: [128]
     device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=360 max=360 min=360 avg=360
 13: data region reached 2 times
  17: data copyout transfers: 1
   device time(us): total=11 max=11 min=11 avg=11
```

module.f90

```fortran
4   module var
5     integer :: i
6     integer, parameter :: maxi=10000
7     integer :: a(maxi)
8     !$ACC declare copyout( i, maxi, a(:) )
9   end module var
10
11  program loop
12    use var
13    !$ACC parallel loop
14    do i=1,maxi
15      a(i) = i
16    enddo
17    write(*,*) a(maxi)
18  end program loop
```

78 / 89

Data Region

Time line

Data Region

Important Notes

- Data transfers between the host and the device are costly;
- It is mandatory to minimize these transfers to achieve good performance;
- It is possible to use data clauses whitin **kernels**/**parallel** and/or **data** regions;
   - The **update** directive can be used to avoid unexpected bahaviors.

Asynchronism

Introduction

- By default kernels are executed synchronously;

- The accelerator is able to manage several execution threads, running concurrently;

- In order to achieve better performance it is recommended to maximize overlaps between:
  - Computations and data transfers;
  - kernel/kernel if they are independent.

Asynchronism is activated by adding the async(execution thread number) clause to one of these directives: **parallel, kernels, serial, enter data, exit data, update** and **wait**.

In all cases **async** is optional.

It is possible to specify a number inside the clause to create several execution threads.
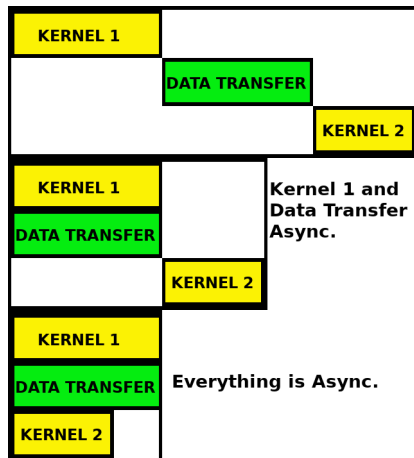
Asynchronism

sync.f90 − s=10000

```
25    !$ACC enter data create( a(1:s), b(1:s),&
26    !$ACC& c(1:s) )
27    ! b is initialized on host
28    do i=1,s
29      b(i) = i
30    enddo
31    !$ACC parallel loop
32    do i=1,s
33      a(i) = 42
34    enddo
35    ! Update vector b located on device
36    ! with data from the host
37    !$ACC update device(b)
38    !$ACC parallel loop
39    do i=1,s
40      c(i) = 1
41    enddo
42    !$ACC exit data delete( a(1:s), b(1:s),&
43    !$ACC& c(1:s) )
```

```
async:
25, Generating enter data create(c(1:s),b(1:s),a(1:s))
31, Generating Tesla code
  32, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
31, Generating implicit copyout(a(1:10000)) [if not already present]
37, Generating update device(b(:))
38, Generating Tesla code
  39, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
38, Generating implicit copyout(c(1:10000)) [if not already present]
42, Generating exit data delete(c(1:s),b(1:s),a(1:s))
```

```
async  NVIDIA  devicenum=0
time(us): 26
25: data region reached 1 time
31: compute region reached 1 time
  31: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
  elapsed time(us): total=357 max=357 min=357 avg=357
31: data region reached 2 times
37: update directive reached 1 time
  37: data copyin transfers: 1
    device time(us): total=19 max=19 min=19 avg=19
38: compute region reached 1 time
  38: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=3 max=3 min=3 avg=3
  elapsed time(us): total=30 max=30 min=30 avg=30
38: data region reached 2 times
42: data region reached 1 time
```

Asynchronism

async-1.f90 – s=10000

```
25    !$ACC enter data create( a(1:s), b(1:s),&
26    !$ACC& c(1:s) )
27    ! b is initialized on host
28    do i=1,s
29      b(i) = i
30    enddo
31    !$ACC parallel loop async(1)
32    do i=1,s
33      a(i) = 42
34    enddo
35    ! Update vector b located on device
36    ! with data from the host
37    !$ACC update device(b)
38    !$ACC parallel loop
39    do i=1,s
40      c(i) = 1
41    enddo
42    !$ACC exit data delete( a(1:s), b(1:s),&
43    !$ACC& c(1:s) )
```

```
async:
25, Generating enter data create(c(1:s),b(1:s),a(1:s))
31, Generating Tesla code
  32, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
31, Generating implicit copyout(a(1:10000)) [if not already present]
37, Generating update device(b(:))
38, Generating Tesla code
  39, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
38, Generating implicit copyout(c(1:10000)) [if not already present]
42, Generating exit data delete(c(1:s),b(1:s),a(1:s))
```

```
async  NVIDIA  devicenum=0
time(us): 23
25: data region reached 1 time
31: compute region reached 1 time
  31: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=355 max=355 min=355 avg=355
31: data region reached 2 times
37: update directive reached 1 time
  37: data copyin transfers: 1
    device time(us): total=16 max=16 min=16 avg=16
38: compute region reached 1 time
  38: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=3 max=3 min=3 avg=3
    elapsed time(us): total=33 max=33 min=33 avg=33
38: data region reached 2 times
42: data region reached 1 time
```

Asynchronism

async-2.f90 – s=10000

```
25    !$ACC enter data create( a(1:s), b(1:s),&
26    !$ACC& c(1:s) )
27    ! b is initialized on host
28    do i=1,s
29      b(i) = i
30    enddo
31    !$ACC parallel loop async(1)
32    do i=1,s
33      a(i) = 42
34    enddo
35    ! Update vector b located on device
36    ! with data from the host
37    !$ACC update device(b) async(2)
38    !$ACC parallel loop
39    do i=1,s
40      c(i) = 1
41    enddo
42    !$ACC exit data delete( a(1:s), b(1:s),&
43    !$ACC& c(1:s) )
```

```
async:
25. Generating enter data create(c(1:s),b(1:s),a(1:s))
31. Generating Tesla code
  32. !$acc loop gang, vector(128) ! blockidx%x threadidx%x
31. Generating implicit copyout(a(1:10000)) [if not already present]
37. Generating update device(b(:))
38. Generating Tesla code
  39. !$acc loop gang, vector(128) ! blockidx%x threadidx%x
38. Generating implicit copyout(c(1:10000)) [if not already present]
42. Generating exit data delete(c(1:s),b(1:s),a(1:s))
```

```
Timing may be affected by asynchronous behavior
set PGI_ACC_SYNCHRONOUS to 1 to disable async() clauses
async  NVIDIA  devicenum=0
time(us): 25
25: data region reached 1 time
31: compute region reached 1 time
  31: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=378 max=378 min=378 avg=378
31: data region reached 2 times
37: update directive reached 1 time
  37: data copyin transfers: 1
    device time(us): total=17 max=17 min=17 avg=17
38: compute region reached 1 time
  38: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=35 max=35 min=35 avg=35
38: data region reached 2 times
42: data region reached 1 time
```

Asynchronism

async-3.f90 – s=10000

```
25   !$ACC enter data create( a(1:s), b(1:s),&
26   !$ACC& c(1:s) )
27   ! b is initialized on host
28   do i=1,s
29     b(i) = i
30   enddo
31   !$ACC parallel loop async(1)
32   do i=1,s
33     a(i) = 42
34   enddo
35   ! Update vector b located on device
36   ! with data from the host
37   !$ACC update device(b) async(2)
38   !$ACC parallel loop async(3)
39   do i=1,s
40     c(i) = 1
41   enddo
42   !$ACC exit data delete( a(1:s), b(1:s),&
43   !$ACC& c(1:s) )
```

```
async:
25, Generating enter data create(c(1:s),b(1:s),a(1:s))
31, Generating Tesla code
 32, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
31, Generating implicit copyout(a(1:10000)) [if not already present]
37, Generating update device(b(:))
38, Generating Tesla code
 39, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
38, Generating implicit copyout(c(1:10000)) [if not already present]
42, Generating exit data delete(c(1:s),b(1:s),a(1:s))
```

```
Timing may be affected by asynchronous behavior
set PGI_ACC_SYNCHRONOUS to 1 to disable async() clauses
async  NVIDIA  devicenum=0
time(us): 22
25: data region reached 1 time
31: compute region reached 1 time
 31: kernel launched 1 time
  grid: [79]  block: [128]
   device time(us): total=4 max=4 min=4 avg=4
  elapsed time(us): total=369 max=369 min=369 avg=369
31: data region reached 2 times
37: update directive reached 1 time
 37: data copyin transfers: 1
  device time(us): total=15 max=15 min=15 avg=15
38: compute region reached 1 time
 38: kernel launched 1 time
  grid: [79]  block: [128]
   device time(us): total=3 max=3 min=3 avg=3
  elapsed time(us): total=33 max=33 min=33 avg=33
38: data region reached 2 times
42: data region reached 1 time
```

Asynchronism

Wait Clause

- The **wait** clause can be used in of dependent kernels;
- **Wait** can be used alone or with a list of **kernels** numbers, *i.e.* **wait(1,2)**.

Kernel 2 waits Kernel 1 and data transfer

Asynchronism

wait.f90 – s=10000

```
25   !$ACC enter data create( a(1:s), b(1:s),&
26   !$ACC& c(1:s) )
27   ! b is initialized on host
28   do i=1,s
29     b(i) = i
30   enddo
31   !$ACC parallel loop async(1)
32   do i=1,s
33     a(i) = 42
34   enddo
35   ! Update vector b located on device
36   ! with data from the host
37   !$ACC update device(b) async(2)
38   !$ACC parallel loop async(3)
39   do i=1,s
40     c(i) = 1
41   enddo
42   !$ACC parallel loop wait(2)
43   do i=1,s
44     b(i) = b(i)*i
45   enddo
46   !$ACC exit data delete( a(1:s), b(1:s),&
47   !$ACC& c(1:s) )
```

```
async:
25, Generating enter data create(b(1:s),c(1:s),a(1:s))
31, Generating Tesla code
  32, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
31, Generating implicit copyout(a(1:10000)) [if not already present]
37, Generating update device(b(:))
38, Generating Tesla code
  39, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
38, Generating implicit copyout(c(1:10000)) [if not already present]
42, Generating Tesla code
  43, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
42, Generating implicit copy(b(1:10000)) [if not already present]
46, Generating exit data delete(c(1:s),b(1:s),a(1:s))
```

```
Timing may be affected by asynchronous behavior
set PGI_ACC_SYNCHRONOUS to 1 to disable async() clauses
async  NVIDIA  devicenum=0
time(us): 28
25: data region reached 1 time
31: compute region reached 1 time
  31: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=4 max=4 min=4 avg=4
    elapsed time(us): total=382 max=382 min=382 avg=382
31: data region reached 2 times
37: update directive reached 1 time
  37: data copyin transfers: 1
    device time(us): total=16 max=16 min=16 avg=16
38: compute region reached 1 time
  38: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=3 max=3 min=3 avg=3
    elapsed time(us): total=34 max=34 min=34 avg=34
38: data region reached 2 times
42: compute region reached 1 time
  42: kernel launched 1 time
    grid: [79]  block: [128]
    device time(us): total=5 max=5 min=5 avg=5
    elapsed time(us): total=21 max=21 min=21 avg=21
42: data region reached 2 times
46: data region reached 1 time
```

GPU Debugging

PGI Auto-compare for OpenACC

- Results can diverge between programs running on a CPU versus a GPU due to programming errors, precision of numerical intrinsics, or variations in compiler optimizations.
- OpenACC auto-compare runs compute regions redundantly on both the CPU and GPU.
- When data is copied from the GPU back to the CPU, GPU results are compared with those computed on the CPU.
- Auto-compare works on both structured and unstructured data regions, with difference reports controlled by environment variables to quickly pinpoint where results start to diverge and adapt the program or compiler options as needed.

GPU Debugging

PGI Auto-compare for OpenACC

The auto-compare is activated during compilation as the race condition example, **pb-sync.f90**:

- **pgf90 -acc -ta=tesla:cc35,cc60,autocompare -Minfo=accel pb-sync.f90**

```
23      !$acc parallel
24      !$acc loop gang
25      do i=1,nx
26        a(i) = 1.0_8
27      enddo
28      !$acc loop gang reduction(+:summ)
29      do i=nx,1,-1
30        summ = summ + a(i)
31      enddo
32      !$acc end parallel
```

```
reduction:
23, Generating Tesla code
  25, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
  29, !$acc loop gang, vector(128) ! blockidx%x threadidx%x
  Generating reduction(+:summ)
23, Generating implicit copy(summ) [if not already present]
  Generating implicit copyout(a(:nx)) [if not already present]
```

```
PCAST Float summ in function reduction, pb-sync.f90:32
idx: 0 FAIL ABS  act: 5.00267200e+07 exp: 1.67772160e+07 dif: 3.32495040e+07
compared 2 blocks, 100000001 elements, 400000004 bytes
1 errors found in 1 blocks
absolute tolerance = 0.00000000000000000e+00, abs=0

Accelerator Kernel Timing data
pb-sync.f90
reduction  NVIDIA  devicenum=0
time(us): 41,027
23: compute region reached 1 time
  23: kernel launched 1 time
    grid: [65535]  block: [128]
    device time(us): total=9,036 max=9,036 min=9,036 avg=9,036
    elapsed time(us): total=9,077 max=9,077 min=9,077 avg=9,077
  23: reduction kernel launched 1 time
    grid: [1]  block: [256]
    device time(us): total=83 max=83 min=83 avg=83
    elapsed time(us): total=111 max=111 min=111 avg=111
23: data region reached 2 times
  23: data copyin transfers: 1
    device time(us): total=7 max=7 min=7 avg=7
  32: data copyout transfers: 25
    device time(us): total=31,901 max=1,433 min=12 avg=1,276
```